SmartD: Smart Meter Data Analytics Dashboard*

Aylin Jarrah Nezhad, Tri Kurniawan Wijaya, Matteo Vasirani, and Karl Aberer
School of Computer and Communication Sciences
École Polytechnique Fédérale de Lausanne (EPFL)
CH-1015 Lausanne, Switzerland
{aylin.jarrahnezhad, tri-kurniawan.wijaya, matteo.vasirani, karl.aberer}@epfl.ch

ABSTRACT

The ability of smart meters to communicate energy consumption data in (near) real-time enables data analytics for novel applications, such as pervasive demand response, personalized energy feedback, outage management, and theft detection. Smart meter data are characterized by big volume and big velocity, which make processing and analysis very challenging from a computational point of view. In this paper we presented SmartD, a dashboard that enables the data analyst to visualize smart meter data and estimate the typical load profile of new consumers according to different contexts, temporal aggregations and consumer segments.

Categories and Subject Descriptors

H.2.8 [**Database Applications**]: Data mining; G.3 [**Probability and Statistics**]: Time series analysis

Keywords

smart meters; visualization; energy consumption analysis

1. INTRODUCTION

Future ICT-based energy systems will rely on an Advanced Metering Infrastructure (AMI), a system that measures and collects data about energy usage and power quality using smart meters installed at the consumer premises [5]. Smart meter data has an important role in several Smart Grid applications and enables novel data analytics tasks, such as energy consumption behavior analysis, theft detection, outage management, pervasive demand response at residential level, and personalized energy feedback. However, processing and analyzing smart meters data is very challenging, because it is characterized by big volume and big velocity, and how to *extract* useful information from it is still an open question.[1]

In this paper we present SmartD, a dashboard for smart meter data visualization and analysis. SmartD has been built to be (i) seamlessly integrated with existing data collection infrastructures, (ii) intuitive to use, and (iii) easy to extend. To visualize and analyze smart meter data, SmartD supports context selection (e.g., summer, winter, weekend, or weekdays), different temporal aggregations (e.g., hourly, daily, or weekly), and consumer selection (either individual or clusters of consumers). Although this functionality is commonly found in other energy dashboard or time series visualization, SmartD's additional and novel contributions are: (i) estimating the typical hourly load profiles based on demographic information, and (ii) determining the attributes of the demographic profile that are relevant to consumer's energy consumption behavior for a given context. This functionality can be used to predict the typical load profile of new consumers, or to understand the energy consumption behavior of different consumers (e.g., employed vs retired, family vs single).[2]

2. SMARTD

We developed SmartD on top of GSN [2], a widely used middleware for sensor networks deployment. Given that smart meters are essentially sensors, GSN can be seamlessly integrated with an existing smart metering infrastructure, enabling applications running on GSN to receive real-time smart meter readings (push mode) as well as obtaining them from a DB or text files (pull mode). Figure 1 shows the architecture diagram of SmartD.

SmartD needs to be able to (i) retrieve and process smart meter data with big volume and velocity, and (ii) visualize and extract valuable information from that data. While the first capability is provided by GSN, we briefly explain the second in the following sections. We remark that although for demonstration purposes we use the Irish CER dataset [1][3], SmartD can be used with any time series of smart meter data and consumer demographic profiles in the form of $<attribute,value>$ tuples.

2.1 Energy Consumption Analysis

For the visualization of energy consumption data (see Figure 2), SmartD has several key features, detailed below.

Temporal aggregations. SmartD supports different time granularities, from half-hourly to monthly. In addition, a set of basic statistical aggregation functions is also provided, such as sum, average, min, and max.[4]
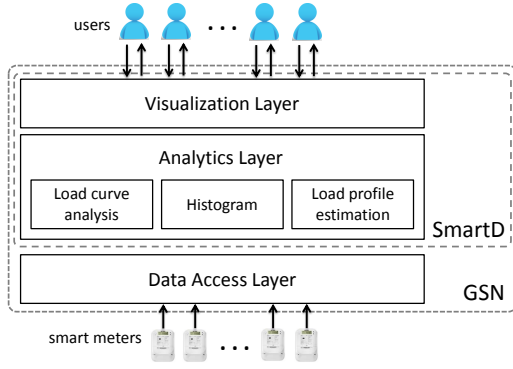
Figure 1: SmartD architecture diagram

Figure 2: Energy consumption analysis

Consumer aggregations. SmartD supports visualization of energy consumption of a set of individual consumers, clusters of consumers, or a mix of both. We use a simple but flexible grammar $id\left((,id) \mid (;id)\right)^*$ to specify the desired visualization, where id is a consumer identifier, the character ";" separates clusters, and the character "," separates consumers within a cluster. An individual consumer is then expressed as a cluster of one consumer. If a cluster of more than one consumer is specified, the users can choose the functions to aggregate the energy data within the cluster (such as sum, average, min, or max).

Consumer characterization. SmartD provides an option to focus more on consumer demand shape, by plotting z-normalized data. This functionality can be used, e.g., to spot consumers who have morning peak, evening peak, or both.

Histogram. SmartD provides a histogram view with the distribution of energy consumption values, which can be useful to analyze the way people consume energy. For example, we found that the energy consumption of residential consumers follows a log-normal distribution, peaked around their base load, while commercial and industrial consumers follows a normal distribution, peaked around the mean consumption of working hours.

2.2 Energy Consumption Estimation

SmartD supports data analysts by providing insights related to energy consumption behavior. First, it helps to answer questions about consumer load profile given her demographics, such as *what is the difference between load profile of families with and without children?*, or *can we estimate the typical load profile of a new consumer using her socio-demographic information?*. SmartD estimates the load profile of a consumer (see Figure 3), if provided with (a subset of) the consumer demographic information, as well as the context of interest (e.g., weekend, Monday, summer, etc.).
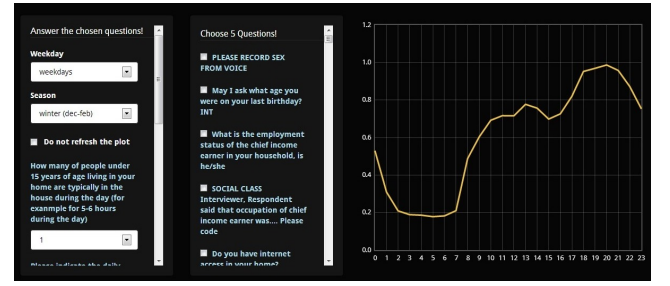

Figure 3: Energy consumption estimation

This is done using *k-nearest neighbor* (*k-NN*), where the best k is determined using *leave-one-out cross-validation*.

Second, SmartD also helps to infer demographic information which significantly influence energy consumption on a specific context, e.g., weekend, Monday, or summer. This is implemented using correlation-based feature selection [4] and *k-NN*, where the features are the demographic information and the target classes are the hourly consumption values. These learning functionalities is developed using the WEKA machine learning library [3].

3. CONCLUSION

In this paper, we presented SmartD, a dashboard for smart meter data analysis and visualization developed on top of GSN. SmartD has been released as an open-source project.[5] SmartD capabilities include visualization of energy consumption data and estimation of the typical load profile of a consumer according to her demographic and contextual information. As future extensions, other functionalities such as customer segmentation [7], interpolation of missing values, and load forecasting [6], could be added.

4. REFERENCES

[1] Smart Metering Trial Data Publication. The Commission for Energy Regulation (CER), 2012.
[2] K. Aberer, M. Hauswirth, and A. Salehi. A Middleware for Fast and Flexible Sensor Network Deployment. In *Proceedings of the 32nd International Conference on Very Large Data Bases*. VLDB Endowment, 2006.
[3] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten. The WEKA Data Mining Software: An Update. *SIGKDD Explor. Newsl.*, 11(1):10–18, Nov. 2009.
[4] M. A. Hall. *Correlation-based Feature Subset Selection for Machine Learning*. PhD thesis, University of Waikato, Hamilton, New Zealand, 1998.
[5] D. Hart. Using AMI to Realize the Smart Grid. In *Power and Energy Society General Meeting - Conversion and Delivery of Electrical Energy in the 21st Century, 2008 IEEE*, pages 1–2, 2008.
[6] S. Humeau, T. K. Wijaya, M. Vasirani, and K. Aberer. Electricity Load Forecasting for Residential Customers: Exploiting Aggregation and Correlation between Households. In *Sustainable Internet and ICT for Sustainability (SustainIT), 2013*, pages 1–6, Oct 2013.
[7] T. K. Wijaya, T. Ganu, D. Chakraborty, K. Aberer, and D. P. Seetharam. Consumer Segmentation and Knowledge Extraction from Smart Meter and Survey Data. In *SIAM International Conference on Data Mining (SDM14)*, 2014.

[5] See SmartD's source code, demo video, and supplementary material for this paper at https://github.com/LSIR/smartd

(Supplementary Material)
SmartD: Smart Meter Data Analytics Dashboard [*]

Aylin Jarrah Nezhad, Tri Kurniawan Wijaya, Matteo Vasirani, and Karl Aberer

School of Computer and Communication Sciences

École Polytechnique Fédérale de Lausanne (EPFL)

CH-1015 Lausanne, Switzerland

{aylin.jarrahnezhad, tri-kurniawan.wijaya, matteo.vasirani, karl.aberer}@epfl.ch

1. SUPPLEMENT FOR SECTION 2.1

Consumer aggregations.

A simple example of a grammar which expresses consumer selection/aggregation: string "$1; 2; 3, 4, 5$" corresponds to the visualization of the energy data of consumer 1 alone, consumer 2 alone, and a cluster composed of consumers 3, 4, and 5.

2. SUPPLEMENT FOR SECTION 2.2

Estimating typical load profile.[1]

SmartD is able to estimate consumer typical load profile given her demographics and contextual information. Let D be the set of demographic information, and C be the context that we are interested in. The set of demographic information, D, can be, for example: a family with two children, live in 2000 sq. ft. apartment, and own a dishwasher. The context, C, can be, for example: weekdays in January, or Monday in the summer. In addition, let N be the set of $k \in \mathbb{N}$ consumers with the closest demographics to D. Thus, $|N| = k$. Then, the estimated load profile of consumers with demographics D on context C is the average of (hourly) load profile of consumers in N.

A question remains, however, to decide the best k. Should k be 1, 2, 3, or something else? To answer this, for each k under consideration, we perform *leave-one-out-cross-validation*. See Algorithm 2.1 for details. For load profiles L_i and L_j, function $dist(L_i, L_j)$ return the distance between L_i and L_j. It can be computed, for example, using the difference between the norm of L_i and L_j.

Discovering significant demographic characteristics.

SmartD is also able to infer demographic information that significantly influences energy consumption on a specific context, e.g., weekdays in January, or Monday in the summer. For this purpose, we use a supervised feature selection algorithm, namely *correlation-based feature selection*.[2] We refer to this algorithm as *cfs*.

Let an *instance* be a tuple (F, l), where $F = \{f_1, \ldots, f_{|F|}\}$ is a feature set and l is a target attribute. Given a set of in-

---

[1] We use the terms *energy consumption* and *load* interchangeably.

[2] See the bibliographic information for this method in the main paper.

Algorithm 2.1: Find the best k

Input: a set of consumers $\mathcal{A}$, a set of k under consideration $\mathcal{K} = \{k_1, \ldots k_n\}$, contextual information C

Output: the best $k \in \mathcal{K}$

1 **foreach** $k \in \mathcal{K}$ **do**
2    $\delta_k \leftarrow 0$
3   **foreach** $i \in \mathcal{A}$ **do**
4      $\mathcal{A}' \leftarrow \mathcal{A} \setminus i$
5   Let N be the set of k consumers in $\mathcal{A}'$ having the closest demographics to i
6      $L_i \leftarrow$ (hourly) load profile of i
7      $L_N \leftarrow$ average (hourly) load profile of consumers in N on context C
8      $\delta_k \leftarrow \delta_k + dist(L_i, L_N)$

9 **return** $\arg\min_k(\delta_k)$

---

stances $\mathcal{I}$, applying *cfs* to $\mathcal{I}$ results in the set R of indexes of the features that are deemed to be relevant to the target attributes. Formally $cfs(\mathcal{I}) = R = \{r_1, \ldots, r_{|R|}\}$ such that $r_i \subseteq \{1, \ldots, |F|\}$.

Next, we explain how to infer top-q demographic characteristics which are relevant to the energy consumption for a context C. Let $D = \{d_1, \ldots, d_{|D|}\}$ be the set of consumer demographics. We define F_i as the feature set of consumer i, where each of its element is consumer i's demographic information. Thus $|F| = |D|$. Let l_i^h be the average of hourly energy consumption of consumer i, on context C, at hour $1 \leq h \leq 24$. Further, let $\mathcal{A}$ be the set of consumers, and $\mathcal{I}^h$ be the set of instances for hour h, consist of tuples (F_i, l_i^h) for all consumers $i \in \mathcal{A}$.

For $1 \leq h \leq 24$, let $cfs(\mathcal{I}^h) = R^h$. Then, we define $score(r) = |\{R^h \mid r \in R^h, 1 \leq h \leq 24\}|$, for $1 \leq r \leq |F|$. The top-q demographic characteristics of the set of consumers $\mathcal{A}$ on context C is the q demographics $d_{r_1^*}, \ldots, d_{r_q^*}$ with the highest scores. That is, the top-q demographics are $d_{r_1^*}, \ldots, d_{r_q^*}$, where $score(r^*) \geq score(r)$ for all $r^* \in \{r_1^*, \ldots, r_q^*\}$ and $r \in \{1, \ldots, |F|\} \setminus \{r_1^*, \ldots, r_q^*\}$.

Summary of Revision

Paper title: "SmartD: Smart Meter Data Analytics Dashboard"
a demo paper, accepted in e-Energy 2014, Cambridge, U.K.

April 16, 2014

General Remarks

We would like to thank the anonymous reviewers for their helpful comments. What we meant by the *github page* is `https://github.com/LSIR/smartd`.

Response to Reviewer 1

1.1 **Comment:** Critical components missing - how is this different from other energy dashboards proposed in research and existing in industry

Response: SmartD provides (1) data visualization and (2) analytics to estimate the load profile and infer demographic attributes that are relevant to the energy consumption. While the first functionality is commonly found in other energy dashboard or time series visualization tools, the second is SmartDs novel contribution. We have added this information to the revised paper as well.

1.2 **Comment:** Critical components missing - some benchmarking for their dashboard (how does it scale for large volume/missing datasets etc.) which are common characteristics for smart meter data.

Response: SmartD is developed on top of GSN. Its capability to receive and process data is inherited from GSN as well. For GSN performance we refer readers to:

> K. Aberer, M. Hauswirth, and A. Salehi. A Middleware for Fast and Flexible Sensor Network Deployment. In Proceedings of the 32nd International Conference on Very Large Data Bases. VLDB Endowment, 2006.

What SmartD adds are the functionality to *plot* (visualize) the data, and the *analytics layer*. Both of them require shorter processing time compared to the time needed to retrieve the data. However, we agree that benchmarking the visualization and analytics functionality is needed. We plan to perform such evaluation in the future and post it in the github page.

1.3 **Comment:** I have not seen any smart metering deployment collecting data using GSN - can authors point to any such existing deployment and hence motivate why a visualization over GSN will be useful

Response: GSN has been widely used in sensor network deployments. The list of deployments, among others, includes: geosciences, hydrology, health monitoring, and pollution. See, e.g., `http://sourceforge.net/apps/trac/gsn` and `http://sourceforge.net/apps/trac/gsn/wiki/other_use_cases`. In our case, a smart meter is a sensor that measures energy consumption, thus it can be well handled by GSN. What we want to emphasize here is the *versatility* of GSN. More specifically, GSN is able to retrieve data, not only from the sensor deployment, but also from the utility company database (even from *frequently updated* csv files). Developed on top of GSN, SmartD benefits from this versatility as well. Whether the current deployments use GSN or not, SmartD is still relevant:

- if the deployment uses GSN, then there is no problem,
- if the deployment does not use GSN, then SmartD is still able to retrieve the data from the database server, either directly or indirectly (for example, using csv files).

1.4 **Comment:** While you have motivated the paper describing high volume and velocity nature of the data, it is not clear how you solved these issues in your implementation. This may be worthy of explaining if the paper gets accepted.

Response: We agree that SmartD needs to (i) retrieve and process smart meter data with big volume and velocity, and (ii) visualize and extract valuable information from that data. While the first capability is provided by GSN, we explain the second in Section 2.1 and 2.2 in the main paper.

1.5 **Comment:** It may be further useful to give some numbers about dataset being used to evaluate the dashboard and provide some microbenchmarks for the same.

Response: The Irish CER dataset contains half-hourly measurements of around 5,000 customers over 1.5 years. The customers consist of residential houses and small and medium-sized enterprises. The measurements started in July 2009 and ended in December 2010. Since the trial was about dynamic pricing, we used only the data from the control group, composed of customers who are not affected by the different pricing schemes. More specifically, we chose residential customers that belong to the control group and have no missing values. This results in the selection of 782 customers. Because of the space limitation, we added only a brief information about the dataset in the main paper. Additionally, we are not allowed to provide the dataset in the github repository, since an agreement with CER has to be signed in order to obtain this data. See `http://www.ucd.ie/issda/data/commissionforenergyregulationcer/` for more information.

1.6 **Comment:** It is also mentioned that SmartD is easy to extend. However, no such information about APIs is provided even on the github link provided in the paper that may substantiate this comment.

Response: We have now added information about this to the github page.

Response to Reviewer 2

2.1 **Comment:** The tool seems to be built exactly for the purpose of understanding the specific CER data set.

Response: SmartD can be used with any time series of smart meter data and consumer demographic profiles in the form of $<attribute, value>$ tuples. We have now added this information as well to the main paper. The use of CER dataset in the main paper (including Figure 2 and 3) is only for demonstration purpose.

2.2 **Comment:** It is unclear how SmartD relates to other existing solutions for time series visualization.

Response: See Response 1.1.

2.3 **Comment:** The authors name smart meter analysis challenging due to "size" and "velocity" of the data, it is not clear from the paper how the tool addresses these challenges as the data set used for evaluation is relatively small and static (i.e., no stream data).

Response: See Response 1.4.

2.4 **Comment:** The technical details in section 2.2 are difficult to understand as they are too condensed and hardly motivated.

Response: We have revised Section 2.2 in the main paper. In addition, we added a more detailed explanation about this in the supplementary material that can be found in the github page. We hope that these revision made the paper clearer.

2.5 **Comment:** The authors present SmartD as a "first step for an open, advanced, and useful smart meter data analytics dashboard." However they do not say what functionality they would include into such a dashboard.

Response: We have added the future work to the main paper.

2.6 **Comment:** The analysis presented in the paper should be given in more detail: Why is a classification approach needed? The paper says "This functionality is based on k-NN classification and correlation-based feature selection [4], where the features are the demographic characteristics and the target classes

are the hourly consumption values of consumers for the given context." I do not understand how (and why) electricity consumption - a continuous property - can be a target class for a k-NN classification. This needs more explanation in the paper. Why not select a subset of properties (that are apparently selected by their influence on the electricity consumption) and create the mean value of the consumption of all households that have the same value for the selected properties?

Response: The last sentence of this comment illustrates what actually SmartD does to *estimate consumer typical load profile* given certain demographic information. We have revised Section 2.2 in the main paper to make it more clear. In addition, we provided more detailed explanation about this aspect in the supplementary material that can be found in the github page.

2.7 **Comment:** Two typos: though → through; ".." → "."

Response: We have fixed the typos.

Response to Reviewer 3

3.1 **Comment:** The paper describes a visualization and analysis tool for smart meter data. The interesting aspect of the tool is that it can provide contextual load estimations depending on demographic information. However, there is no discussion on related visualization tools

Response: See Response 1.1.

3.2 **Comment:** The last paragraph in Sec 2.2 is very vague. While I understand that a reasonable description of the classification mechanism cannot be given in 7-8 lines, the authors may consider to renounce giving an incomplete and confusing description and rather insist on other novel features of SmartD, especially concerning the future ones.

Response:

- See Response 2.4 for more details about the classification (learning) mechanism.
- SmartD's novel features have been given an additional emphasize in the Section 1 in the main paper.
- We have now mentioned several possible interesting extensions for SmartD in the Section 3 in the main paper.